

1 Introduction

- The questions we ask of each database:
 - *What type of DB is this?* Relational etc.
 - *What was the motivation?* Who developed and why.
 - *How do you talk to it?* Through shell, programming, protocols?
 - *Why is it unique?*
 - Performance/scalability.

- Large genres of DBs:

Relational Two-dimensional tables from set theory, queries are in Structured Query Language (SQL), based on relational set theory. PostgreSQL.

Key-Value Literally a hash. Riak + Redis.

Columnar Instead of storing rows of a table together, store columns together. Easy to build sparse attributes. HBase.

Document Stores hashes of basically anything (JSON). MongoDB + CouchDB.

Graph Stores nodes and relationships between nodes, can traverse along relationships effectively. Neo4J.

- Best practice is obviously to use multiple DBs for diff use cases.

Glossary:

Relational Based on relational algebra, not based on relations between tables.

CRUD Create Read Update Delete. Every other operation is a higher order composition of these.

REST REpresentational State Transfer, guideline for mapping CRUD resources to URLs.

MapReduce Not a new concept, but succinctly captured by *It is faster to send the algorithm to the data than the data to the algorithm.* (Note: Generally best when reducing in MapReduce to reduce to the same schema as mapped values to allow chaining reduces.)

CAP Theorem When a network partition error occurs and messages are lost in a *partitioned* network, you can only guarantee *consistency* (across partitions) or *availability* to further incoming requests. More precisely, *at any given moment in time you cannot be consistent, available and partition tolerant.*

2 PostgreSQL

- Roots in 1970s, supported SQL by 1996. Relational. Archetypally stable/reliable.
- Relations = TABLES, attributes = COLUMNS, tuples = ROWS.
- On INSERT, can specify RETURNING to get any automatically populated values e.g. SERIAL primary keys.
- Joins

Inner Join Join two columns from two tables on equality of those columns.

Outer Join Join two columns from two tables and for at least one of the two tables always return even if the lookup in the other table fails.

- *Indexing* helps avoid full table scans. PostgreSQL automatically indexes the primary key and all `UNIQUE` attributes. Can do either hash or btree indexes. Also when specify `FOREIGN KEY`, index target table.
- Can `INSERT INTO` values that are `SELECTED` from another table! Handy to prevent hardcoding primary keys everywhere.
- Aggregate functions allow post-processing, e.g. `count()`. Can `GROUP BY` aggregate functions and can also filter on aggregated values with `HAVING` the same way `SELECT` filters with `WHERE`.
- Can use `PARTITION BY` to not collapse rows within each group. Use case is when `SELECTING` over a field not used in an aggregate query and so can get conflicting values when also `GROUP BY`ing.
- Transactions + ACID compliance:
 - Transactions ensure that every command of a set is executed else none.
 - ACID—Atomic (all or nothing), Consistent (never stuck in inconsistent state e.g. nonexistent foreign keys), Isolated (transactions do not interfere), Durable (committed transactions will always endure even if server crashes).
- Can store procedures (`FUNCTIONS`) that are loaded by the database side. Obviously faster than postprocessing returned data from the db but higher maintenance cost.
- Can specify `TRIGGERS` that hit these stored procedures
- Lifecycle of a SQL command: parsed into query tree, modify used off rules (and views, which are a specific type of rule), hit query planner, executed and return.
- Can specify custom rules e.g. how to interpret certain operations on a view.
- Can fuzzy string search using `LIKE` and `ILIKE`, can regex or even Levenstein (edit distance), trigram. All have corresponding indexes that can be built, all pluggable using PostgreSQL-exclusive packages.
- The `cube` seems neat, you can define feature vectors for each row and tell PostgreSQL how to measure distances between feature vectors and query on said distance.
- Apparently does not scale well horizontally b/c partitioning is difficult for relational databases. But is very good for normalized data, extremely reliable w/ transactions + ACID compliance.

3 Riak

- Riak is a highly distributed, highly available key-value DB that is built for a web interface, notably to be `cURL`ed.
- All nodes are equal, very easy to join nodes `riak-admin join`.
- Insert by `curl -X PUT`, delete by `curl -X DELETE`.
- Key format is `<server>/<bucket>/<key>` and the key can be auto-generated on insertion.

- Can `Link` values to metadata labels. Say that the label `contains` the value. Is one-way pointer.
- Can query on links, called *link walking*, with `GET` to `<server>/<link bucket>/<link>`. Can specify to *keep* each step in the link walking.
- Can also tag with various other metadata, `X-Riak-Meta-*`. Can also specify MIME type to store images etc. that are linked to existing entries.
- Can execute commands by posting JSON bodies to endpoints e.g. `/mapred`. Endpoints often take a function body in plaintext, so can generally point to a bucket + key instead. Stored procedures!
- Riak supports filtering keys prior to map reduce.
- The *Riak Ring* is key to consistency and durability;
 - All nodes are peers, growing and shrinking the cluster is trivial.
 - Riak uses a 160-bit *ring* hashing keys to determine which Riak servers store the values for which key.
 - Riak then partitions the ring and each server claims partitions sequentially on startup.
 - Riak accomplishes redundancy by hashing each key to `N` nodes, considering a write successful when `W` nodes have completed the write, and then you can specify reading from `R` nodes.
- Writes in Riak are by default not durable and not written to disk before acknowledgement.
- Riak by default 204s when writing to a server that is not yet up, and a neighboring node will buffer the write! Careful for *cascading failure* when the neighboring nodes fail consecutively due to overload.
- Riak handles concurrent writes by effectively tagging each update w/ a commit history, and conflict resolution must be performed manually by specifying which `Vclocks` an update overwrites.
 - The commit history, called the *vector clock*, is pruned as more updates to a value occur, configurable per bucket.
- Can specify pre/post commit hooks, notably validators and/or computed values.
- Search with Apache Solr interface, `/solr`.

4 HBase

- Apache HBase is made for very big data, on the order of GBs (EN:?? That's really small.).