

## Blend in Chicago: MongoDB World 2017

Yubo Su

## Blend

June 21, 2017

- Tom Schenk, Chief data officer, Chicago. *WindyGrid*.
  - Track colocated data, 911 calls to Tweets to weather.
  - Flexible schema: {what, when, where}
  - Predictive analytics (example, where to send food inspector) using visualization of multiple causal layers.
- Dev Ittycheria, CEO MongoDB
  - 2007 is watershed year, AWS, iPhone, Android, and many others.
  - Argue b/c storage costs dropped below a critical point.
  - MongoDB also in 2007: document model, distributed systems + aggregation.

- Eliot Horowitz, CTO, MongoDB
  - 3.6 ships November, already on Github.
  - MongoDB Charts (3.6)
    - Business Intelligence: BI Connector is SQL interface.
    - Coercing data to table is difficult: polymorphic schemas, arrays.
    - Solution: *MongoDB Charts!* Data visualization tool, handles above.
  - 3.6 document model features:
    - `$lookup` takes sub-pipelines!
    - `$update` can operate on arrays natively! Takes a filter over array entries, can iterate over nested.
    - JSON Schemas.
  - 3.6 distributed systems:
    - Native retryable writes
    - *Change Streams* can get a stream of changes to a db.

- Eliot Horowitz, CTO, MongoDB (continued)
  - Mongo Atlas
    - “Should be irresponsible to run MongoDB in cloud w/o Atlas”
    - Built in security, one-click spin up, built in scaling elasticity.
    - Data browser + performance viewer in UI (utilization stats, examine queries as stream, explore data),
    - Live migration service (not very live in demo, requires downtime for mirror to catch up and change source of truth).
    - Now with MS Azure + Google Cloud support too (+ AWS).
  - Performance Adviser.
  - CRUD support in data browser.
  - Charts!
  - LDAP Auth.
  - Cross-region, cross-cloud!
- MongoDB Stitch (Beta as of today in Atlas, 06/20/17)
  - “Backend as a service”
  - REST API for MongoDB
  - Configuration-based auth/security
  - Service composition to govern how services talk to each other.

# Squeezing the Most out of Your Document Model

06/20/17 1050-1130: Norberto Leite, Lead Curriculum Engineer, MongoDB

- Nested schema, spectrum of highly normalized or denormed storage.
  - Normalized requires foreign keys, requires looking into many collections.
  - Denorm is simpler query, complex schema.
- Consider three possible behaviors:
  - Get player: Denorm outperforms.
  - Add new field to doc: either add new collection or modify every doc, the same.
  - Change existing field: If a highly shared field, normalized is very fast.
- Optimizing highly normalized:
  - Can optimize with aggregate, but more importantly `db.createView()`.
    - Views are basically stored aggregates.
    - Better `$project` support.
  - Also consider, if reading much more than writing, should store calculated fields!
- Optimizing denormed:
  - Should normalize fields that are infrequently updated.
- `tl;dr` normalized have fast write, slow reads. Should embed everything that is infrequently updated.

# Advanced Schema Design Patterns

06/20/17 1140–1220: Daniel Coupal, Senior Curriculum Engineer, MongoDB

- Axiom: data models maximize performance + scalability despite latency, costs, hardware.
- Common issues #1, too many optional fields:
  - Use attribute array, `[{key: keyName, value}]`.
  - Accommodates optional fields.
- Common issues #2, working set does not fit in RAM.
  - Can subset, truncate data
  - Probably also useful for showing users too.
- Common issues #3, data consistency.
  - Accept instantaneous inconsistency, duplicate at regular intervals ☺.
- Common issues #4, repeated computations
  - Reads generally outnumber writes, apply computation on write.
- Common issues #5, expensive tracking
  - e.g. expensive to increment on every page view
  - Solution: random number in range  $[1, N]$ , increment by  $N$ .
- Common issues #6, large data easily overflow
  - Bucket, store buckets into a separate collection.

- Microservices vs. monolith, preferable b/c web scale, faster iteration, compartmentalized.
- One common rule of thumb is that one developer can own the whole thing, a couple hundred lines, but not everybody
- Hard metal vs. Docker (Kubernetes) vs. Atlas.
- Kafka can run general events while Mongo streams (the new feature) only handles database updates.

# Index Usage for Nested Logical Queries

06/20/17 1440-1520: Tess Avitabile, Software engineer, MongoDB


- Query system overview:
  - Input: JSON
  - Parse into tree
  - Generate plan (which indices for which leaves of the tree)
  - Plan selection: try all of the plans for a trial period, see which one was fastest (Note: plan caching)
  - Execution & return
- ORs inside of ANDs is a pain for plan generation.
  - AND is considered indexed when one child has index.
  - OR is considered indexed when all children have indices.
  - ORs have to dedupe by hashing to merge the two results.
- Problem: no tight index bounds on these queries
  - *Tight* index bounds are when all documents in index bounds match the query.
  - (As opposed to when a parent node imposes a filter, `FETCH`)
- Bounds are not tight b/c two branches of children cannot talk to each other! e.g. the OR will not be tight since the AND above will have to further fetch against its other child.



# Index Usage for Nested Logical Queries (cont)

06/20/17 1440-1520: Tess Avitabile, Software engineer, MongoDB

- Solution: Disjunctive Normal Form?
  - AND with OR child solved!
  - Exponentially many plans though, index choices at each child.
- Solution: OR-pushdowns! Predicates pulled up to the AND parent and pushed down into any OR children if they can tighten index bounds.
  - Note that this is not imposed as an extra AND condition, just metadata for the recursive query planner to plan against.
- Paper: Query Optimization by Predicate Move-Around



**Figure:** AND with OR child. Consider if index is  $\{f1, f2\}$ ?  $\{f2, f1\}$ ?

# Multi-Master architectures in MongoDB

06/20/17 1530–1610: Pavel DUCHOVNY, TSE, MongoDB

- Key to geographic colocation.
- Zone sharding + replica sets
  - Zone sharding: shards per region.
  - Replica sets are `mongod` processes that share the same data.
- Configuration is:
  - One primary in each region, each a separate zone shard.
  - One secondary in its own region (prio 3), two in another (prio 2), and two more in a third (prio 1, 0) for symmetry across all regions, hidden secondary.
  - Spread across multiple regions, odd number of voting members, primary DC members should have higher priorities.
- Can specify region on read/write.
- Upshot is that can do multi-region writes while guaranteeing local availability on read.
- Configurable to write to secondary especially if primary lives in a different region.
- Can configure with `MAX_STALENESS` parameter for when a cluster can be read from.

# Globally Distributed RESTful Object Storage

06/20/17 1620-1700: Julio Viera, Backend VP, Fuze

- Built an object store for internal communications, chat + attachment retention, Mongo backbone.
- RESTful so easy to expose HTTP link as a db.
- Nested schema corresponding to URL:  
`/users/:id/chat/convs/X/messages/Y`
- Storage (chat), collection (convs), sub-collection (messages), documentIds
- Pubsub (user is online) can be done by consuming the oplog on the db primary.
- To shard and hide it to the user, just need some lookups `userId` → sharding keys.

- Saska Mojsilovic, IBM
  - Need for more data in health for precision health service distribution.
  - All sorts of orgs estimating and predicting from sparse data.
- Claudio Gosiker (Florida Blue) & Alan Chhabra (MongoDB)
  - Use data for healthcare outreach, personalizable views for customer reps.
- Matt Parker, Stand-up Mathematician!

- Bjorn Freeman-Benson, CTO, InVisionApp
  - Via microservices, can stand up new cluster in 10m!
  - Also has a `bailey stage it` etc.
  - QA against EA customers, automatically rolls out to rest of customers afterwards (24h).
- Cisco moved eCommerce to MongoDB, 40b connections?
- Justin Moses, Lead Software Engineer, MongoDB
  - Data auralization vs. visualization!
  - `npmjs link`
  - Just turns numbers into music.
- Jane McGonigal, Game Designer, AvantGame
  - 2.1b gamers > 1h/day, more stats.
  - 72% workforce not engaged, vs. 80+% of schoolchildren engaged.
  - Consider: *"Opposite of play is not work but depression."*
  - Video games overstimulate brain regions exactly what depression suppresses.
  - Pokemon Go fitness lol.
  - Reality's obligation to engage the way video games do, AR > VR!

# Migrating from EC2 to Atlas

06/21/17 1050–1130: Jesse Dearing, SRE, InVision

- Mongo at InVision
  - 28 replica sets 4 env
  - 2000 rps, 600 wps
  - Chef to manage EC2, Mongo
- Old stack:
  - EC2 instance, deploy, manually configure/shard
  - Manual: backups, monitoring, alerting, security, updates
- Atlas:
  - All above, REST API, dashboards
- Transition Preparation
  - SSL (Atlas mandatory)
  - AWS VPC Peering
  - VPN + security setup, Amazon DNS
  - MongoDB 3.x + WiredTiger
- Transition
  - UI Live Migrator (<1m downtime for oplog)
  - `mongomirror` for full ZDT: Initial sync, streams oplog, point to new instance *before* fully synced, continues syncing.
  - Full ZDT but momentary inconsistency (< 1s).
  - In case of rolling ZDT deploys when re-pointing, inconsistency is most noticable; graceful degradation!
- Epilogue
  - Alerts, new playbooks, backup restores.
  - Automatic provisioning for new services that need MongoDB.

**MONGODB**  
**WORLD'17**

blend

# Design Patterns Resilient to Infrastructure Failure

06/21/17 1140-1220: Feng Qu, Senior MTS, eBay

- Availability =  $\frac{MTTF}{MTTF + MTTR}$ ,  
mean time to failure/recovery.
- High write (?) 3+2+2, 2 slaves each  
except maybe 2 arbiters (only used for  
quorum voting).
  - 2 slaves on master so if one slave  
fails, master still has a slave
  - Master fail → slave
  - Primary datacenter fail → re-elect in  
second datacenter
- High read goes to 3+3+3
- Shard for overall read/write scaling.
- For latency, `readPreference=nearest` or  
`readPreference=secondaryPreferred`.
- `WriteConcern majority`.
- Tagged reads: read only from nodes  
that made it into the write majority, get  
updated data w/o paying remote read  
penalty.
- tl;dr many patterns given different  
read/write patterns latency vs.  
throughput questions.
- Mongo supports causal consistency.

# Write and Read Concern

06/21/17 1400–1440: Alex Komyagin, Senior Consulting Engineer, MongoDB

- How Mongo writes (3-replica)
  - App sends doc to primary, primary writes to in-memory structures and oplog, sends ack back to app.
  - In no particular order: journal, replicated to secondaries, then eventually (~ seconds) data files.
- How writes disappear
  - What if primary fails right before replication + data files?
  - One secondary is promoted, primary reboots.
  - Once primary reboots as secondary, produces rollback containing in-transit write `bson`.
  - App cannot find in-transit write!
- Many people disable rollback ☹☹☹
- RCA: Ack before replicated.
- Solution: `writeConcern`, # replicas or `majority`.
- Rec: majority for important writes, else 2 (lower latency, but at least one other server knows).
- Dirty reads: detect when read is of unstable state. Modes:
  - Local
  - Majority
  - Linearizable (blocks until last write is replicated)
- Linearizable can block e.g. example above, when master goes down, since the write never gets replicated.

**MONGODB**  
**WORLD'17**

blend



# Write and Read Concern

06/21/17 1400–1440: Alex Komyagin, Senior Consulting Engineer, MongoDB

- Majority vs. Linearizable
  - Majority returns the majority that a particular node knows about.
  - Linearizable guarantees recency, but can only work on primary, very slow.
- Retryable writes can handle write concern timeout.

# Common Cluster Configuration Pitfalls

06/21/17 1450–1530: Alex Komyagin, Technical Services Engineer, MongoDB

- Replication improves availability.
- Arbiters vote but store no data.
- Consider a consumer that needs high availability, inexpensive “report generation.”
  - Say one data center has primary + arbiter, another has secondary,
  - If first data center crashes, the secondary only has 1/3 votes, cannot become new primary!
  - This is so if there is an LB, the secondary (which only sees itself) can never see the primary + arbiter, so the secondary cannot promote itself to primary just b/c primary + arbiter down, cannot differentiate between down vs. LB.
- Propositions continued:
  - Put the arbiter in a separate instance?
  - If the primary goes down, then cannot writeConcern majority!
- Solution: 1 primary 2 secondaries.
- Suppose the primary has super high priority.
  - Recall that Mongo is eventually consistent, so secondaries can lag badly if under provisioned.
  - If primary goes down, comes back up as a secondary, rollback, ☹.
  - Solution: Equally provision secondaries!
- All writes go to primary.
- Can configure secondary reads, but Mongo is eventually consistent so dangerous.

**MONGODB**  
**WORLD'17**

blend

# Common Cluster Configuration Pitfalls

06/21/17 1450–1530: Alex Komyagin, Technical Services Engineer, MongoDB

- Sharding: replica set no bandwidth.
  - For scalability.
  - Divides *collection* across replica sets.
  - **Send shard key with query, else queries all shards!**
  - Sharding works by chunking the shard key.
  - If a chunk grows to over 20% of max size, `mongos` tries to split the chunk. If a key range cannot be split, it “lacks cardinality.”
  - Chunk splitting is slow. Primaries know what data ranges they own but secondaries do not, so during copying of splitting can have duplicate docs.
- Chunk balancer only balances based on number of shards, not chunk size!
- Be sure write concern can be fulfilled if one node crashes.
- No arbiters.
- Design replica sets for availability, then add.
- Shard key selection! Should have a wide range of different values but recall must be included in query!
- Secondary reads are discouraged, especially on sharded.
- Best practices for zone sharding (tag shards with zones) are similar.

# MongoDB Performance in Theory and Practice

06/21/17 1540–1620: Baron Schwartz, Founder/CEO, VividCortex

- Performance is either user latency or resource consumption.
- Residence time for a request = queue + service time.
- Profiling is key! Identify obviously bad queries by sorting by time.
- Another trick is to sort by frequency, repetitive calls are also awful.
- Do not use MongoDB's slow query log!
- `db.serverStatus()`, `currentOp()` have lots of useful info and are built in, but should not poll these.

# Powerful Analysis with the Aggregation Pipeline

06/21/17 1630–1710: Asya Kamsky, Lead PM, MongoDB

- `tl;dr` can do data analysis in-db, faster!
- Use `{explain:true}` to get an explain of the plan that Mongo is using to evaluate the aggregate, since Mongo does a lot of fancy push-arounds.
- `$let` seems to be good for readability.
- `$map`, `$reduce`, `$filter` for arrays!
- Refactor into helper functions that reduce stages of the aggregate pipeline.