

## Blend in Chicago: MongoDB World 2017

Yubo Su

## Blend

June 21, 2017

- Tom Schenk, Chief data officer, Chicago. *WindyGrid*.
  - Track colocated data, 911 calls to Tweets to weather.
  - Flexible schema: {what, when, where}
  - Predictive analytics (example, where to send food inspector) using visualization of multiple causal layers.
- Dev Ittycheria, CEO MongoDB
  - 2007 is watershed year, AWS, iPhone, Android, and many others.
  - Argue b/c storage costs dropped below a critical point.
  - MongoDB also in 2007: document model, distributed systems + aggregation.

- Eliot Horowitz, CTO, MongoDB
  - 3.6 ships November, already on Github.
  - MongoDB Charts (3.6)
    - Business Intelligence: BI Connector is SQL interface.
    - Coercing data to table is difficult: polymorphic schemas, arrays.
    - Solution: *MongoDB Charts!* Data visualization tool, handles above.
  - 3.6 document model features:
    - `$lookup` takes sub-pipelines!
    - `$update` can operate on arrays natively! Takes a filter over array entries, can iterate over nested.
    - JSON Schemas.
  - 3.6 distributed systems:
    - Native retryable writes
    - *Change Streams* can get a stream of changes to a db.

- Eliot Horowitz, CTO, MongoDB (continued)
  - Mongo Atlas
    - “Should be irresponsible to run MongoDB in cloud w/o Atlas”
    - Built in security, one-click spin up, built in scaling elasticity.
    - Data browser + performance viewer in UI (utilization stats, examine queries as stream, explore data),
    - Live migration service (not very live in demo, requires downtime for mirror to catch up and change source of truth).
    - Now with MS Azure + Google Cloud support too (+ AWS).
  - Performance Adviser.
  - CRUD support in data browser.
  - Charts!
  - LDAP Auth.
  - Cross-region, cross-cloud!
- MongoDB Stitch (Beta as of today in Atlas, 06/20/17)
  - “Backend as a service”
  - REST API for MongoDB
  - Configuration-based auth/security
  - Service composition to govern how services talk to each other.

# Squeezing the Most out of Your Document Model

06/20/17 1050-1130: Norberto Leite, Lead Curriculum Engineer, MongoDB

- Nested schema, spectrum of highly normalized or denormed storage.
  - Normalized requires foreign keys, requires looking into many collections.
  - Denorm is simpler query, complex schema.
- Consider three possible behaviors:
  - Get player: Denorm outperforms.
  - Add new field to doc: either add new collection or modify every doc, the same.
  - Change existing field: If a highly shared field, normalized is very fast.
- Optimizing highly normalized:
  - Can optimize with aggregate, but more importantly `db.createView()`.
    - Views are basically stored aggregates.
    - Better `$project` support.
  - Also consider, if reading much more than writing, should store calculated fields!
- Optimizing denormed:
  - Should normalize fields that are infrequently updated.
- `tl;dr` normalized have fast write, slow reads. Should embed everything that is infrequently updated.

# Advanced Schema Design Patterns

06/20/17 1140–1220: Daniel Coupal, Senior Curriculum Engineer, MongoDB

- Axiom: data models maximize performance + scalability despite latency, costs, hardware.
- Common issues #1, too many optional fields:
  - Use attribute array, `[{key: keyName, value}]`.
  - Accommodates optional fields.
- Common issues #2, working set does not fit in RAM.
  - Can subset, truncate data
  - Probably also useful for showing users too.
- Common issues #3, data consistency.
  - Accept instantaneous inconsistency, duplicate at regular intervals ☺.
- Common issues #4, repeated computations
  - Reads generally outnumber writes, apply computation on write.
- Common issues #5, expensive tracking
  - e.g. expensive to increment on every page view
  - Solution: random number in range  $[1, N]$ , increment by  $N$ .
- Common issues #6, large data easily overflow
  - Bucket, store buckets into a separate collection.

- Microservices vs. monolith, preferable b/c web scale, faster iteration, compartmentalized.
- One common rule of thumb is that one developer can own the whole thing, a couple hundred lines, but not everybody
- Hard metal vs. Docker (Kubernetes) vs. Atlas.
- Kafka can run general events while Mongo streams (the new feature) only handles database updates.

# Index Usage for Nested Logical Queries

06/20/17 1440-1520: Tess Avitabile, Software engineer, MongoDB

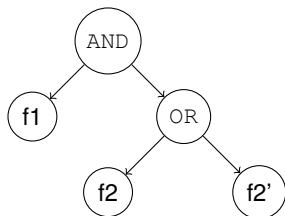
- Query system overview:
  - Input: JSON
  - Parse into tree
  - Generate plan (which indices for which leaves of the tree)
  - Plan selection: try all of the plans for a trial period, see which one was fastest (Note: plan caching)
  - Execution & return
- ORs inside of ANDs is a pain for plan generation.
  - AND is considered indexed when one child has index.
  - OR is considered indexed when all children have indices.
  - ORs have to dedupe by hashing to merge the two results.
- Problem: no tight index bounds on these queries
  - *Tight* index bounds are when all documents in index bounds match the query.
  - (As opposed to when a parent node imposes a filter, `FETCH`)
- Bounds are not tight b/c two branches of children cannot talk to each other! e.g. the OR will not be tight since the AND above will have to further fetch against its other child.



# Index Usage for Nested Logical Queries (cont)

06/20/17 1440–1520: Tess Avitabile, Software engineer, MongoDB

- Solution: Disjunctive Normal Form?
  - AND with OR child solved!
  - Exponentially many plans though, index choices at each child.
- Solution: OR-pushdowns! Predicates pulled up to the AND parent and pushed down into any OR children if they can tighten index bounds.
  - Note that this is not imposed as an extra AND condition, just metadata for the recursive query planner to plan against.
- Paper: Query Optimization by Predicate Move-Around



**Figure:** AND with OR child. Consider if index is  $\{f1, f2\}$ ?  $\{f2, f1\}$ ?

# Multi-Master architectures in MongoDB

06/20/17 1530–1610: Pavel DUCHOVNY, TSE, MongoDB

- Key to geographic colocation.
- Zone sharding + replica sets
  - Zone sharding: shards per region.
  - Replica sets are `mongod` processes that share the same data.
- Configuration is:
  - One primary in each region, each a separate zone shard.
  - One secondary in its own region (prio 3), two in another (prio 2), and two more in a third (prio 1, 0) for symmetry across all regions, hidden secondary.
  - Spread across multiple regions, odd number of voting members, primary DC members should have higher priorities.
- Can specify region on read/write.
- Upshot is that can do multi-region writes while guaranteeing local availability on read.
- Configurable to write to secondary especially if primary lives in a different region.
- Can configure with `MAX_STALENESS` parameter for when a cluster can be read from.

# Globally Distributed RESTful Object Storage

06/20/17 1620-1700: Julio Viera, Backend VP, Fuze

- Built an object store for internal communications, chat + attachment retention, Mongo backbone.
- RESTful so easy to expose HTTP link as a db.
- Nested schema corresponding to URL:  
`/users/:id/chat/convs/X/messages/Y`
- Storage (chat), collection (convs), sub-collection (messages), documentIds
- Pubsub (user is online) can be done by consuming the oplog on the db primary.
- To shard and hide it to the user, just need some lookups userId → sharding keys.

- Saska Mojsilovic, IBM
  - Need for more data in health for precision health service distribution.
  - All sorts of orgs estimating and predicting from sparse data.
- Claudio Gosiker (Florida Blue) & Alan Chhabra (MongoDB)
  - Use data for healthcare outreach, personalizable views for customer reps.
- Matt Parker, Stand-up Mathematician!

- Bjorn Freeman-Benson, CTO, InVisionApp
  - Via microservices, can stand up new cluster in 10m!
  - Also has a `bailey stage it` etc.
  - QA against EA customers, automatically rolls out to rest of customers afterwards (24h).
- Cisco moved eCommerce to MongoDB, 40b connections?
- Justin Moses, Lead Software Engineer, MongoDB
  - Data auralization vs. visualization!
  - `npmjs link`
  - Just turns numbers into music.
- Jane McGonigal, Game Designer, AvantGame
  - 2.1b gamers > 1h/day, more stats.
  - 72% workforce not engaged, vs. 80+% of schoolchildren engaged.
  - Consider: *"Opposite of play is not work but depression."*
  - Video games overstimulate brain regions exactly what depression suppresses.
  - Pokemon Go fitness lol.
  - Reality's obligation to engage the way video games do, AR > VR!

# Migrating from EC2 to Atlas

06/20/17 1050–1130: Jesse Dearing, SRE, InVision

- Mongo at InVision
  - 28 replica sets 4 env
  - 2000 rps, 600 wps
  - Chef to manage EC2, Mongo
- Old stack:
  - EC2 instance, deploy, manually configure/shard
  - Manual: backups, monitoring, alerting, security, updates
- Atlas:
  - All above, REST API, dashboards
- Transition Preparation
  - SSL (Atlas mandatory)
  - AWS VPC Peering
  - VPN + security setup, Amazon DNS
  - MongoDB 3.x + WiredTiger
- Transition
  - UI Live Migrator (<1m downtime for oplog)
  - `mongomirror` for full ZDT: Initial sync, streams oplog, point to new instance *before* fully synced, continues syncing.
  - Full ZDT but momentary inconsistency.
- Epilogue
  - Alerts, new playbooks, backup restores.
  - Automatic provisioning for new services that need MongoDB.

**MONGODB**  
**WORLD'17**

blend